

Solving a PPDDL Problem as a Simple Stochastic Game

Lorenzo Mandelli
1747091

Nicolò Paoletti
1888157

Valerio Ponzi
1760886

Giulio Recupito
1668320

I. Introduction

The aim of this work is to apply the simple stochastic games (SSGs) theory to a classic PPDDL problem, which extends PDDL toward probabilistic scenarios. SSG theory tackle problems as directed graph, considering the possibility of having probabilistic action outcome. Ideally if we had moved from a classic PPDDL to a direct graph of this kind we would have been able to also move from an exponential time to a polynomial one. Initially we will introduce PDDL and PPDDL formally than discussing how these kind of problems are addressed, facing the parsing issue and the graph building. Later we will introduce how we solved the planning problem, introducing SSG theory and the consequent algorithm we end up using. Finally we will conclude with some considerations regarding the computational complexity of our work compared with a classical safe planner.

II. PPDDL Problems

PPDDL stands for "Probabilistic Planning Domain Definition Language", which is an extension of its predecessor: PDDL. This kind of extension is a first step towards a general language for describing probabilistic and decision theoretic planning problems.

i. PDDL

PDDL [3] is a language for specifying deterministic planning domains and problems.

A PDDL planning domain consists of a set T of types, a sub-typing relation $ST \subset T \times T$, a set C of global objects (domain constants), a set P of Predicates, a set F of functions (which in our case are not considered) and a set AS of action schemata.

Predicates are used to encode Boolean state variables. Objects and variables are terms and

each one has a type $\tau \in T$. PDDL also includes support for union types $\tau_1 \cup \dots \cup \tau_n$, with the restriction that each τ_i is a simple type.

Actions in PDDL can be thought like sets of state transitions, with a state being a particular assignment to the set of state variables of a planning problem. An action consists of a precondition, characterizing the set of states in which the action is applicable and an effect, which specifies updates to state variables that occur at the execution in the given state. For a Boolean state variable b , the effect b (or an application yielding the state variable b) simply means that b should be set to true in the next state while, in order to set b to false, the notation (**not** b) is used.

A planning problem consists of a set of state variables V , a set of actions A , an initial state s_0 , a goal condition θ identifying a set of goal states, and an optimization metric f that is typically a function of numeric state variables evaluated in a goal state. In PDDL, a planning problem is always associated with a domain definition, and the definition of a planning problem includes a declaration of a set of problem-specific objects O . The state variables V for the planning problem are obtained from O , C , P , and F as all type-consistent applications of Predicates to objects (including domain constants). The set A of actions is obtained similarly as all type-consistent applications of action schemata in AS to objects in $O \cup C$. The process of obtaining all state variables and actions for a planning problem through the exhaustive application of Predicates and action schemata to objects is referred to as *grounding*.

ii. PPDDL

In order to define probabilistic and decision theoretic planning problems, we need to add support for probabilistic effects. The syntax for

probabilistic effects is

probabilistic $p_1 e_1 .. p_k e_k$

meaning that effect e_i occurs with probability p_i . We require that the constraints $p_i \geq 0$ and $\sum_{i=1}^k p_i = 1$ are fulfilled. However, if the effect is empty, a probability-effect pair can be ruled out. So if the effect (**probabilistic** $p_1 e_1 .. p_k e_k$) with $\sum_{i=1}^l p_i \leq 1$ is the same but easier to read of (**probabilistic** $p_1 e_1 .. p_l e_l q$ (and)) with $q = 1 - \sum_{i=1}^l p_i$.

PPDDL allows arbitrary nesting of conditional and probabilistic effects. A single PPDDL action schema can represent a large number of actions and a single Predicate can represent a large number of state variables. It means that PPDDL often can represent planning problems more synthetically than other representations. For example, the number of actions that can be represented using m objects and n action schemata with arity c is $m \cdot n$, which is not bounded by any polynomial in the size of the original representation $m + n$.

Grounding is not a prerequisite for PPDDL planning, so planners could take advantage of the more compact representation by directly working with action schemata.

iii. Parser

The PPDDL parser we used is an extension of the classical PDDL parser. Our implementation is based on Lex and Yacc. The parser reads the source program and discovers its structure.

Lex and Yacc can generate program fragments that solve this task by decomposing it into sub tasks:

- Split the source file into tokens (Lex);
- Find the hierarchical structure of the program (Yacc);

The first helps to write programs whose control flow is directed by instances of regular expressions in the input stream. Lex source is a table of regular expressions and corresponding program fragments. We extend the set of regular expression in order to include also the possibility of having probability. The newly introduced set of regular expressions handles non-uniform probability density functions along with nested ones. The table is translated to a program which reads an

input stream, copying it into an output stream and partitioning it in strings. The strings have to match the given expressions. As each such string is recognized, the corresponding program fragment is executed. The recognition of the expressions is performed by a deterministic finite automaton generated by Lex. The program fragments written by the user are executed following the order in which the corresponding regular expressions occur in the input stream. The lexical analysis programs written with Lex accept ambiguous specifications and choose the longest match possible at each input point.

On the other hand Yacc provides a general tool for imposing structure on the input to a computer program. Then it generates a parser function to control the input process defined by the user calling the lexical analyzer to pick up the tokens from the input stream.

These tokens are organized according to the input structure rules, called grammar rules. When one of these rules has been recognized, then an action is invoked. The parser produced by Yacc consists of a finite state machine with a stack. The parser is also capable of reading and remembering the next input token (called the lookahead token). The current state is always the one on the top of the stack. The states of the finite state machine are given small integer labels; initially, the machine is in state 0, the stack contains only state 0, and no lookahead token has been read. The machine has only four actions available to it, called shift, reduce, accept, and error.

A move of the parser is done as follows:

1. Based on its current state, the parser decides whether it needs a lookahead token to decide what action should be done; if it needs one, and does not have one, it calls a lexical analyzer function(Lex), representing the kind of token read.
2. Using the current state, and the lookahead token if needed, the parser decides on its next action, and carries it out. This may result in states being pushed onto the stack, or popped off of the stack, and in the lookahead token being processed or left alone.

The reduce action keeps the stack from growing without bounds. Reduce actions are appropriate when the parser has seen the right

hand side of a grammar rule, and is prepared to announce that it has seen an instance of the rule.

The accept action indicates that the entire input has been seen and that it matches the specification. This action appears only when the lookahead token is the end marker, and indicates that the parser has successfully done its job.

The error action, on the other hand, represents a place where the parser can no longer continue parsing according to the specification, hence, when the input tokens seen, together with the lookahead token, cannot be followed by anything that would result in a legal input.

iv. Graph

Based on the aforementioned parser, we can represent an instance of a PPDDL problem as a directed graph. We defined all the classes needed to handle correctly each object composing the PPDDL problem, namely a Domain class, a Problem class and, Literals, Predicates and Terms classes, thus having an object oriented structure underling our programs. Terms are the smaller units needed to build Predicates out of we obtain Literals. For instance considering an example

- `not(vehicle-at('?loc': 'x02y04'))` is an instance of `Literal` class;
- `vehicle-at('?loc': 'x02y04')` is an instance of `Predicate` class;
- `'?loc': 'x02y04'` in an instance of `Term` class.

Actions' preconditions and effects are Predicates, consequently states are represented as set of Predicates. This set is initially taken from the Predicates in the `init` section of the Problem class; from this state we need to compute all possible applicable actions (i.e. match the preconditions with the current state). At this point, we need to distinguish between Predicates without arguments and Predicates with one or more of them: in the first case we just need to check if the precondition is listed in the current state, otherwise the action is not feasible. In the latter case, we need to search all possible Terms that make the action's precondition true, based on the current state's Predicates. This is achieved by querying the current state's Predicates, retrieving all the Term's instances needed by a certain action; then, all

the possible instances of the preconditions are filtered using a natural-join like operation. The action's effect Literals will then be computed from the Terms found in the previous operation. The new state resulting from the current state and a given applicable action is obtained as:

- **Delete** from the current state all the Predicates contained inside the action effects' **negative** Literals
- **Add** in the current state all the Predicates contained inside the action effects' **positive** Literals

If the action is not deterministic (i.e. has a probabilistic effect) the computation of the resulting state is done for each possible outcome of the action. This transitions are stored in a dictionary we called "possible_paths", that has the following structure:

```
{
  "action1(param1, ...)": [state1],
  "action2(param1, ...)": [state2],
  "action3(param1,...)": [state3, state4]
}
```

The graph itself is built recursively. At each iteration we check we are not in the Goal node, nor in a Sink node before computing the "possible paths" dictionary for the current state. Based on the length of this dictionary we understand if we are taking an action with deterministic or probabilistic effects, thus we build accordingly a Max node or an Average node and create the corresponding edges. Each node in the graph contains information about the node type (max, avg, sink and goal) and the node index. Max nodes, differently from Average nodes, contain also the state that the node identifies; Average nodes doesn't contain this information, but will be connected to as much Max nodes as the possible outcome of the action it describes. Edges will contain a representation of the action and the probability of its specific outcome; namely deterministic actions will always connect two Max nodes, with its edges having probability 1, while non deterministic actions will always connect a Max to an Average node with an edge having probability 1, and then connect the Average node with as much edges as the possible outcomes to the resulting Max nodes set. In addition we cannot have edges from Average nodes to Average nodes, since conditional probabilities have

been grounded to a single outcome during the creation of the action. In Figure 1 we show an example of a graph we obtained from a PPDDL problem.




Figure 1: Example of a graph obtained by a simple problem, the red node is the Root and the pink one is the Goal, black nodes are Sink states, in green we coloured Average nodes and in light blue Max nodes

III. Solver

Considering the given obtained representation of our problem, which is a direct graph with three types of vertices respectively, Max, Average, and Sink plus a Start vertex, we are now able to treat it like a simple stochastic game (briefly SSG).

i. SSG

Based on [1] a simple stochastic game is a direct graph $G = (V, E)$. The vertex set V is the union of disjoint sets $V_{Max}, V_{Average}$, together with two special vertices, the 0-Sink and the 1-Sink. One vertex of V is called the *Start* vertex. We assume that each node, different from the Sink and Goal, has one or multiple neighbours, depending on which kind of node it is (Max or Average). We considered an extended version of the initial problem, taking into account the possibility that an Average vertex could have more than two neighbours, which can be associated to a non-uniform probability density function. A *strategy* τ consists of a subset of edges of E : $\tau = (i, j) \subseteq E, \forall i \in V_{Max}$. In the game theory literature, this kind of strategies are called *pure stationary* strategies, since the player do not use probabilistic choices and the

player always chose the same move from a vertex every time the same vertex is reached. In the giving settings it is prove that the SSG can be solved in polynomial time. As sated in [1] the solution of $G = (V, E)$ is the optimal solution of a simple constrained minimization linear programming problem, which is:

$$\begin{aligned}
 \min \quad & \sum_{i=1}^n v(i) \\
 \text{s.t.} \quad & v(i) \geq v(j), \forall i \in V_{Max} \vee (i, j) \in E, \\
 & v(i) \geq \sum_{j=1}^k p_j(v(j)), \forall i \in V_{avg}, k = |neigh(i)|, \\
 & v(i) = 0, \forall i = \text{sink}, \\
 & v(i) = 1, i = \text{goal}, \\
 & v(i) \geq 0, \forall i
 \end{aligned} \tag{1}$$

At this point we know that, having the graph representing such a problem, we can move the task to a simple optimization problem, solvable, for instance by means of the **Simplex Method**.

ii. Simplex Method

In the following lines we will briefly describe how the Simplex method works. The simplex algorithm operates on linear programs in the canonical form

$$\begin{aligned}
 \min_x \quad & c^T x \\
 \text{s.t.} \quad & Ax \leq b, \\
 & x \geq 0
 \end{aligned} \tag{2}$$

In geometric terms, the feasible region defined by all the values of x such that $Ax \leq b$ and $\forall i, x_i \geq 0$ is a (possibly unbounded) convex polytope. An extreme point or vertex of this polytope is known as a basic feasible solution (BFS). The solution of a linear program is accomplished in two steps. In the first step, known as Phase 1, a starting extreme point is found, resulting in a BFS or in an empty region. If the latter is the case then the program is infeasible, if not the simplex algorithm pass to Phase 2, where the algorithm is applied using the BFS found in Phase 1 as a starting point. The possible result from Phase 2 are either an optimum basic feasible solution or an infinite edge on which the objective function is

unbounded above. We can also express a linear program in standard form, so to say when we have a linear program in the form $Ax = b$ $\forall i, x_i \geq 0$ as a tableau:

$$\begin{bmatrix} 1 & -c^T & 0 \\ 0 & A & b \end{bmatrix}$$

The first row defines the objective function and the remaining rows specify the constraints. The zero in the first column represent the zero vector of the same dimension as vector b . If the linear program is in given in the canonical tableau, the simplex algorithm proceeds by performing successive pivot operations each of which give an improved basic feasible solution; the choice of pivot element at each step is largely determined by the requirement that this pivot improves solution.

The simplex method output an array of n elements, one for each node of the graph, containing the probability that starting from that node, we end up in the goal node. Next we display the output of the simplex method applied to the problem depicted in 1

```
[
1.06303393 0.73251663 0.34279100 0.12182052 0.19780839 0.17698863
0.22902882 0.26717919 0.36483302 0.5427862 0.56153425 0.62315621
0.69628452 0.79620168 0.63011896 0.79745855 0.74690075 0.87246654
0.35224907 0.48704109 0.43839591 0.56774975 0.66077537 0.76635253
0.59957841 0.75022523 0.71881281 0.85675294 0.50159683 0.67995671
0.69046894 0.72185974 0.85750786 0.49879535 0.00000000 0.61405444
0.81658037 0.58944646 0.42829789 0.48527385 0.29958237 0.31527231
0.39745978 0.56788973 0.57447170 0.62934697 0.69919646 0.79720542
0.64111624 0.80097843 0.75385140 0.87439825 0.44734644 0.52695710
0.44551900 0.57446107 0.66226406 0.76459594 0.61255951 0.74642802
0.72514508 0.85730757 0.57575163 0.74203877 0.71939429 0.74868643
0.86540956 0.59833805 0.00000000 0.62128621 0.75672717 0.59490724
0.55472726 0.50273828 0.43419941 0.58461589 0.67791959 0.61991466
0.68700579 0.83147315 0.60590508 0.47843292 0.71308211 0.87906198
0.72539078 0.67305875 0.69691321 0.56625850 0.69360311 0.74068538
0.64445998 0.70908954 0.83727118 0.69728012 0.48540678 0.74042209
0.87478646 0.76250429 0.75514588 0.68293597 0.30313191 0.01502747
0.70229479 1. ]
```

The strategy is computed picking from all nodes neighbour to a Max node, the one with higher value, which will more likely lead to a goal node. Next we display the strategy for the same problem depicted in 1

```
0 : strategy: move-car('?to': 'x01y03', '?from': 'x01y01')
2 : strategy: change-tire()
3 : strategy: change-tire()
4 : strategy: move-car('?to': 'x01y03', '?from': 'x02y02')
6 : strategy: load-tire('?loc': 'x03y03')
7 : strategy: change-tire()
8 : strategy: move-car('?to': 'x02y04', '?from': 'x03y03')
10 : strategy: load-tire('?loc': 'x02y04')
11 : strategy: change-tire()
12 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
14 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
16 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
18 : strategy: load-tire('?loc': 'x03y03')
20 : strategy: load-tire('?loc': 'x02y04')
21 : strategy: change-tire()
22 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
24 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
26 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
28 : strategy: move-car('?to': 'x02y04', '?from': 'x03y03')
30 : strategy: change-tire()
31 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
35 : strategy: move-car('?to': 'x01y05', '?from': 'x01y03')
38 : strategy: move-car('?to': 'x01y03', '?from': 'x02y02')
40 : strategy: load-tire('?loc': 'x03y03')
41 : strategy: change-tire()
42 : strategy: move-car('?to': 'x02y04', '?from': 'x03y03')
44 : strategy: load-tire('?loc': 'x02y04')
45 : strategy: change-tire()
46 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
48 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
50 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
52 : strategy: load-tire('?loc': 'x03y03')
54 : strategy: load-tire('?loc': 'x02y04')
55 : strategy: change-tire()
56 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
58 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
60 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
62 : strategy: move-car('?to': 'x02y04', '?from': 'x03y03')
64 : strategy: change-tire()
65 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
69 : strategy: move-car('?to': 'x01y05', '?from': 'x01y03')
72 : strategy: move-car('?to': 'x01y03', '?from': 'x02y02')
74 : strategy: change-tire()
75 : strategy: move-car('?to': 'x02y04', '?from': 'x03y03')
77 : strategy: change-tire()
78 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
81 : strategy: change-tire()
82 : strategy: move-car('?to': 'x01y05', '?from': 'x01y03')
85 : strategy: move-car('?to': 'x01y03', '?from': 'x02y02')
87 : strategy: change-tire()
88 : strategy: move-car('?to': 'x02y04', '?from': 'x03y03')
90 : strategy: change-tire()
91 : strategy: move-car('?to': 'x01y05', '?from': 'x02y04')
94 : strategy: change-tire()
```

Time comparisons		
P01 phase	Our planner	Safe planner
Parsing	0.0030 s	0.0038 s
Graph creation	0.5154 s	/
Solution finding	0.0568 s	/
Total time to solve the problem	0.5727 s	0.01676 s
Complete process time	0.5757 s	0.0207 s
P02 phase	Our planner	Safe planner
Parsing	0.0029 s	0.0022 s
Graph creation	194.8837 s	/
Solution finding	139.5348 s	/
Total time to solve the problem	334.41854 s	0.3543 s
Complete process time	334.4613 s	0.3565 s

Table 1: Comparison between our planner and SP execution time.

```

95 : strategy: move-car('?to': 'x01y05', '?from': 'x01y03')
99 : strategy: move-car('?to': 'x01y03', '?from': 'x01y01')
101 : strategy: load-tire('?loc': 'x02y02')

```

IV. Complexity

In order to analyze the computational complexity of our solution, we have to take into account:

- The computational complexity required for the graph creation
- The computational complexity of the planner

The creation of the graph requires the expansion of each applicable action in each Max node. Thus, the computational complexity is $O(|A|^n)$ where $|A|$ is the set of actions and n is the Maximum distance (depth) from the starting node. As stated in [1], since our SSG involves just Max and Average nodes, the solution of our problem can be found in a time that is polynomial in the number of nodes. In conclusion, as we expected, the graph creation step makes the resulting time complexity exponential.

V. Conclusions

We compared the time results with a very recent planner called Safe-Planner [SP] [2], using two different instances the triangle-tireworld problem, a simpler one [P01] and more complex one, [P02]. As reported in the table 1, SP is exponentially faster in finding a solution, since the most recent planner algorithms

do not expand and create all the search space but exploits different kind of heuristics, unlike our solution that is forced to pass through the graph creation; consequently the simplex method has to work with a number of node which is exponential in the depth of the solution, thus affecting considerably the converging time.

References

- [1] *The Complexity of Stochastic Games* ANNE CONDON, 1992
- [2] *Safe-Planner: A Single-Outcome Replanner for Computing Strong Cyclic Policies in Fully Observable Non-Deterministic Domains* Vahid Mokhtari, Ajay Suresha Sathya, Nikolaos Tsiogkas, Wilm Decre', 2021
- [3] *PDDL - The Planning Domain Definition Language* Ghallab, Malik & Knoblock, Craig & Wilkins, David & Barrett, Anthony & Christianson, Dave & Friedman, Marc & Kwok, Chung & Golden, Keith & Penberthy, Scott & Smith, David & Sun, Ying & Weld, Daniel. (1998).